

1) Problem

Teams waste time and budget searching scattered sources, with no direct bridge from evidence to ex

2) Audience

AI/ML engineers, applied researchers, and product experiment teams under tight delivery and resourc

3) Core Solution (Expert Preflight)

One flow runs five-source similarity search with transparent 0-1 scores, then generates a practical ex

4) Differentiation

This combines scored multi-source retrieval and structured planning from the same prompt context, no

5) Implementation

Node.js + Express onepager with /search and /plan, MCP tool integration, timeout-safe fanout, and pa

6) Reliability & Impact

Best-available outputs continue during partial upstream failures; prompt-hash caching and determinist

GitHub: <https://github.com/hacknation-mintwin/similarity-search>